


Native CPU Attention Thread Sweep

One-layer direct C++ kernel; 32-layer throughput is extrapolated; no-SMT uses taskset without changing system HT

Extrapolated 32-layer CPU-attention throughput



Measured one-layer CPU attention time

