



End-to-End Throughput: GPU-only vs GPU-CPU Heterogeneous



Relative Throughput



Experiment details

Metric: completion_tokens / request_wall_time.

Each workload is run with a fresh vLLM server.

Prefix caching is not disabled, but each server handles only one measured request.

This preserves default vLLM behavior while avoiding repeated-prompt cache hits.

Heterogeneous mode: OMP_NUM_THREADS=64, taskset -c 0-63.

GPU memory utilization: 0.25

Max model length: 10000

Batch size: 1