



End-to-End Throughput: GPU vs. KV-Split CPU Decode

Llama-3.1-8B-Instruct; batch=1; request wall includes prefill and decode



Best KV-split thread count per case

Case	Best	KV tok/s	KV/GPU	Case	Best	KV tok/s	KV/GPU	Case	Best	KV tok/s	KV/GPU
P4K/O100	OMP60	23.24	26.29%	P6K/O100	OMP100	22.23	23.26%	P8K/O100	OMP100	19.37	20.35%
P4K/O300	OMP120	22.75	22.10%	P6K/O300	OMP100	22.35	21.82%	P8K/O300	OMP100	19.59	19.06%
P4K/O400	OMP120	23.31	22.56%	P6K/O400	OMP80	19.13	18.64%	P8K/O400	OMP80	20.69	20.18%