



End-to-End Throughput: GPU vs. No-Split CPU Decode

Llama-3.1-8B-Instruct; batch=1; request wall includes prefill and decode



Best no-split thread count per case

Case	Best	KV tok/s	KV/GPU	Case	Best	KV tok/s	KV/GPU	Case	Best	KV tok/s	KV/GPU
P4K/O100	OMP120	17.13	17.59%	P6K/O100	OMP60	15.68	18.47%	P8K/O100	OMP40	14.02	17.48%
P4K/O300	OMP60	18.26	21.11%	P6K/O300	OMP60	16.17	19.26%	P8K/O300	OMP60	13.93	20.72%
P4K/O400	OMP60	18.03	18.78%	P6K/O400	OMP60	16.32	23.18%	P8K/O400	OMP60	14.35	15.37%