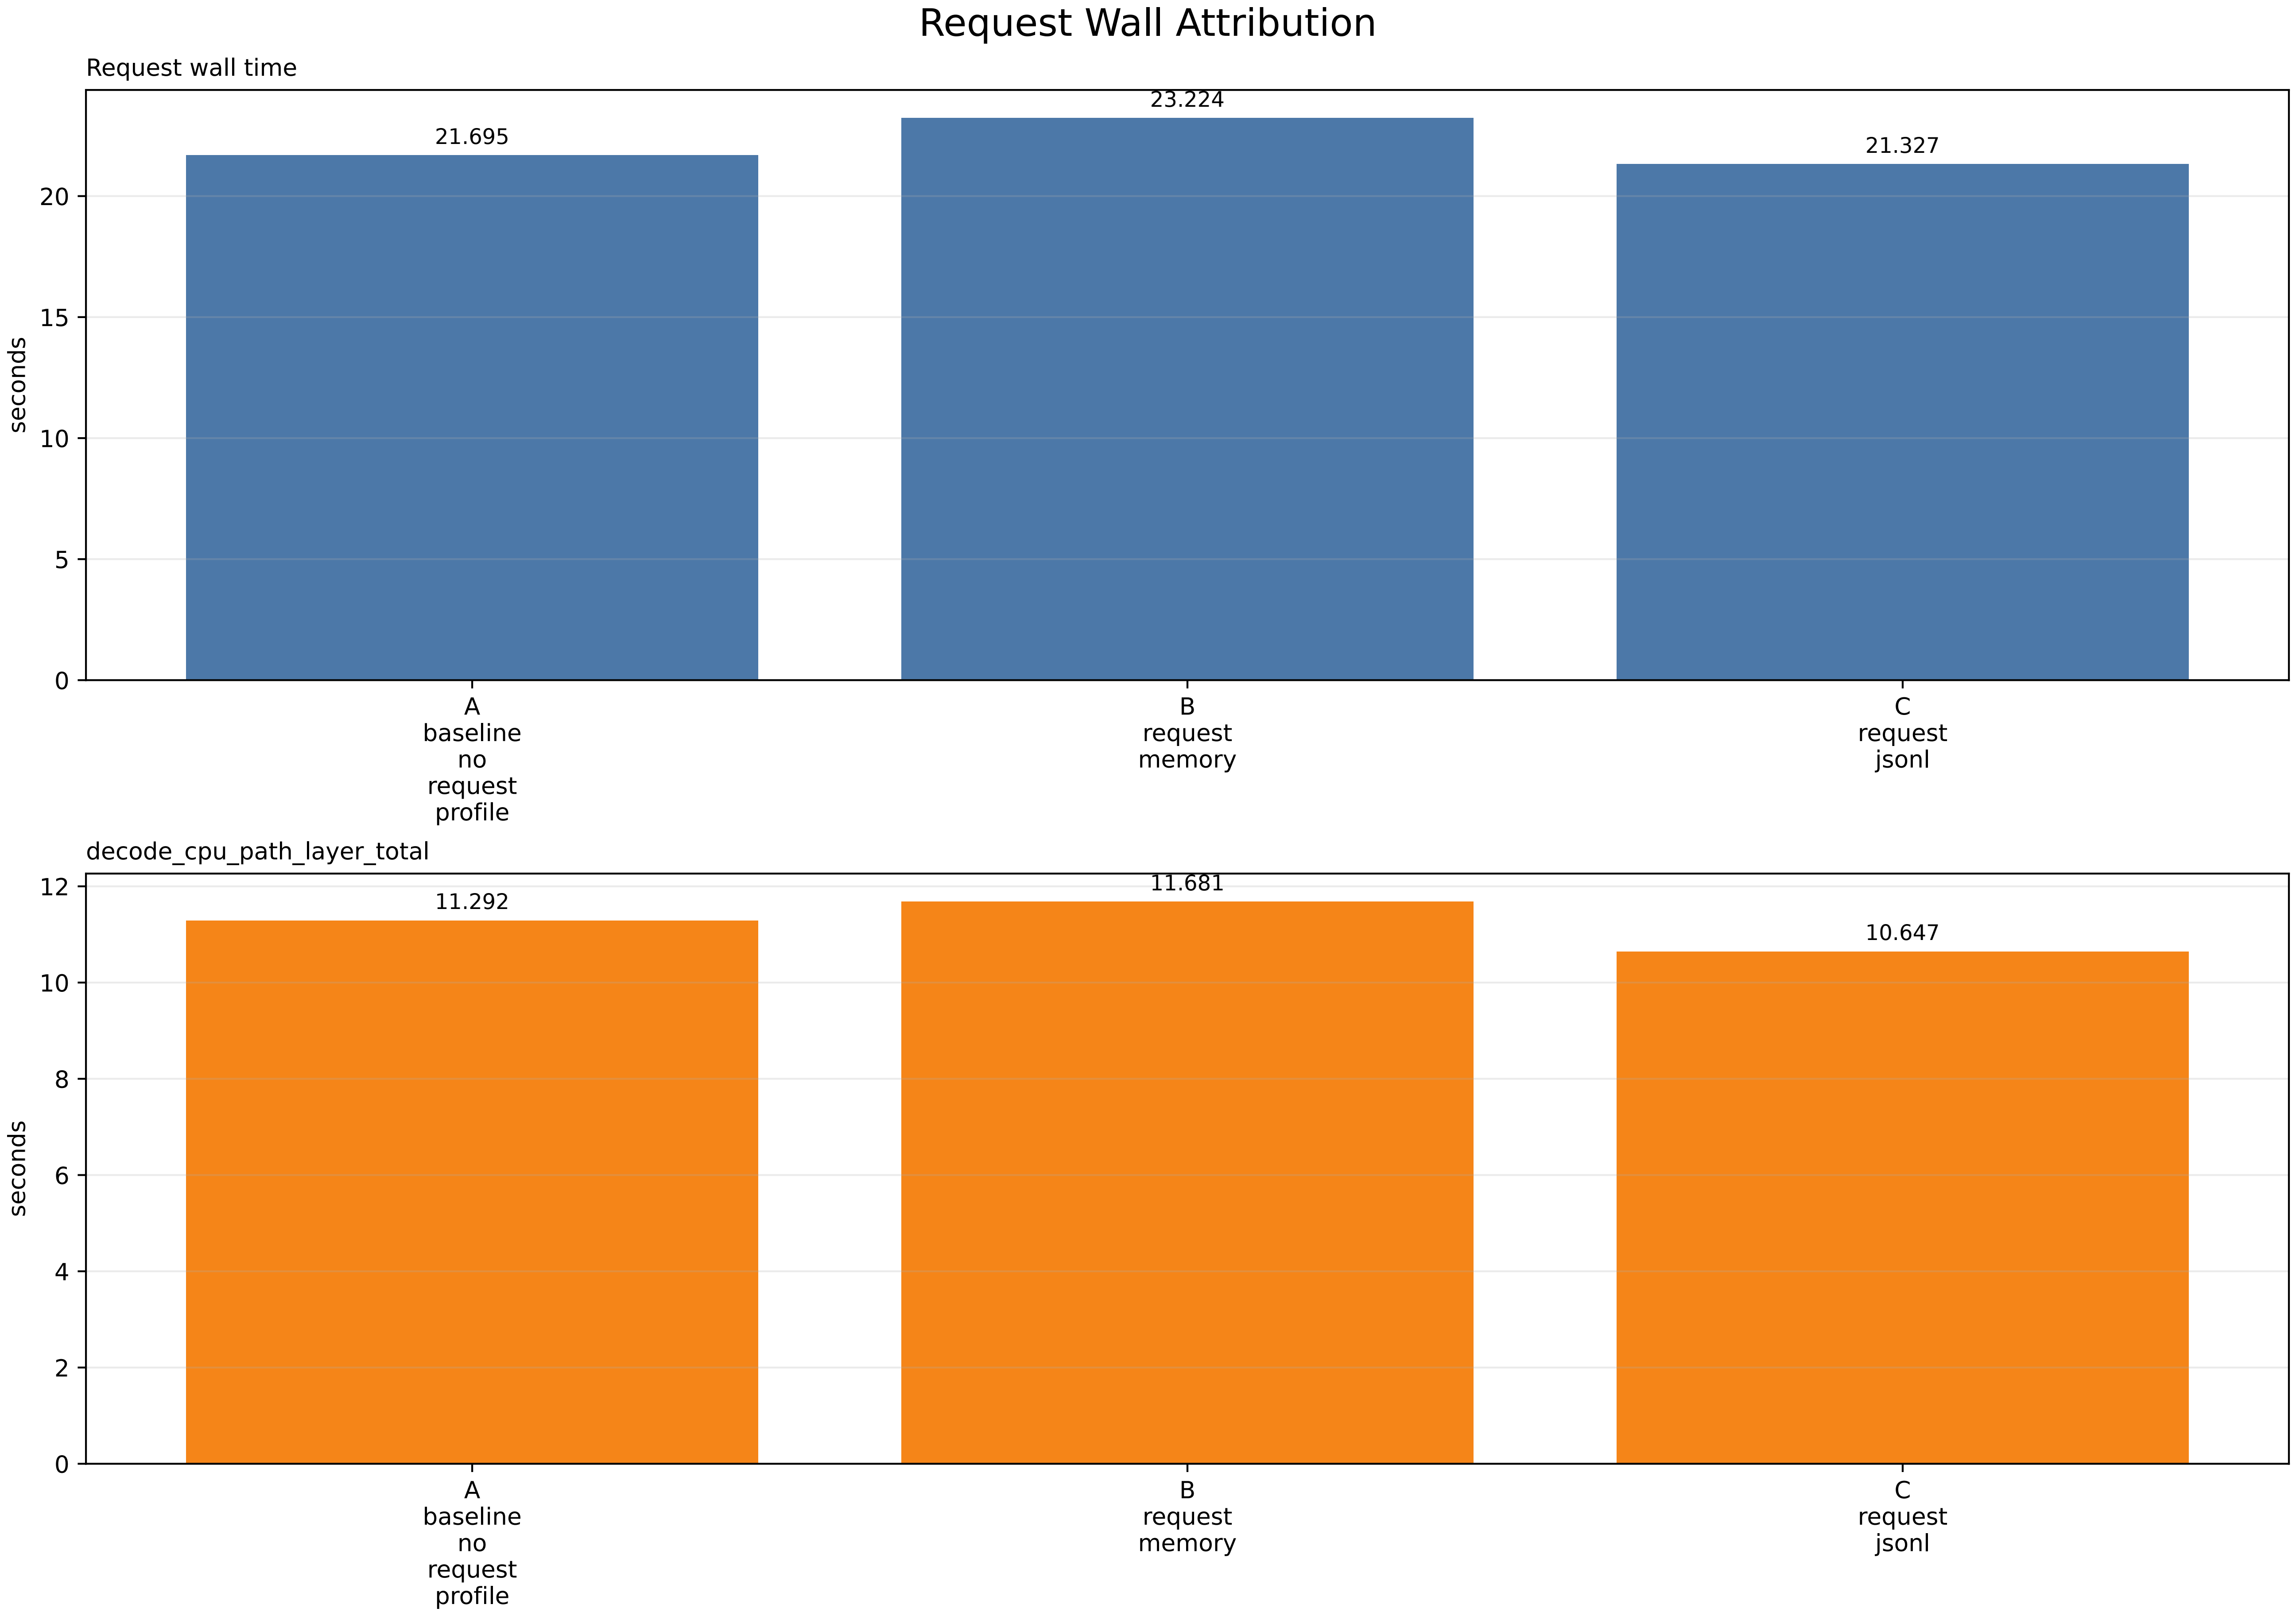
Request Wall Attribution
Request wall time
seconds
0
5
10
15
20
21.695
23.224
21.327
A
baseline
no
request
profile
B
request
memory
C
request
jsonl
decode_cpu_path_layer_total
seconds
0
2
4
6
8
10
12
11.292
11.681
10.647
A
baseline
no
request
profile
B
request
memory
C
request
jsonl

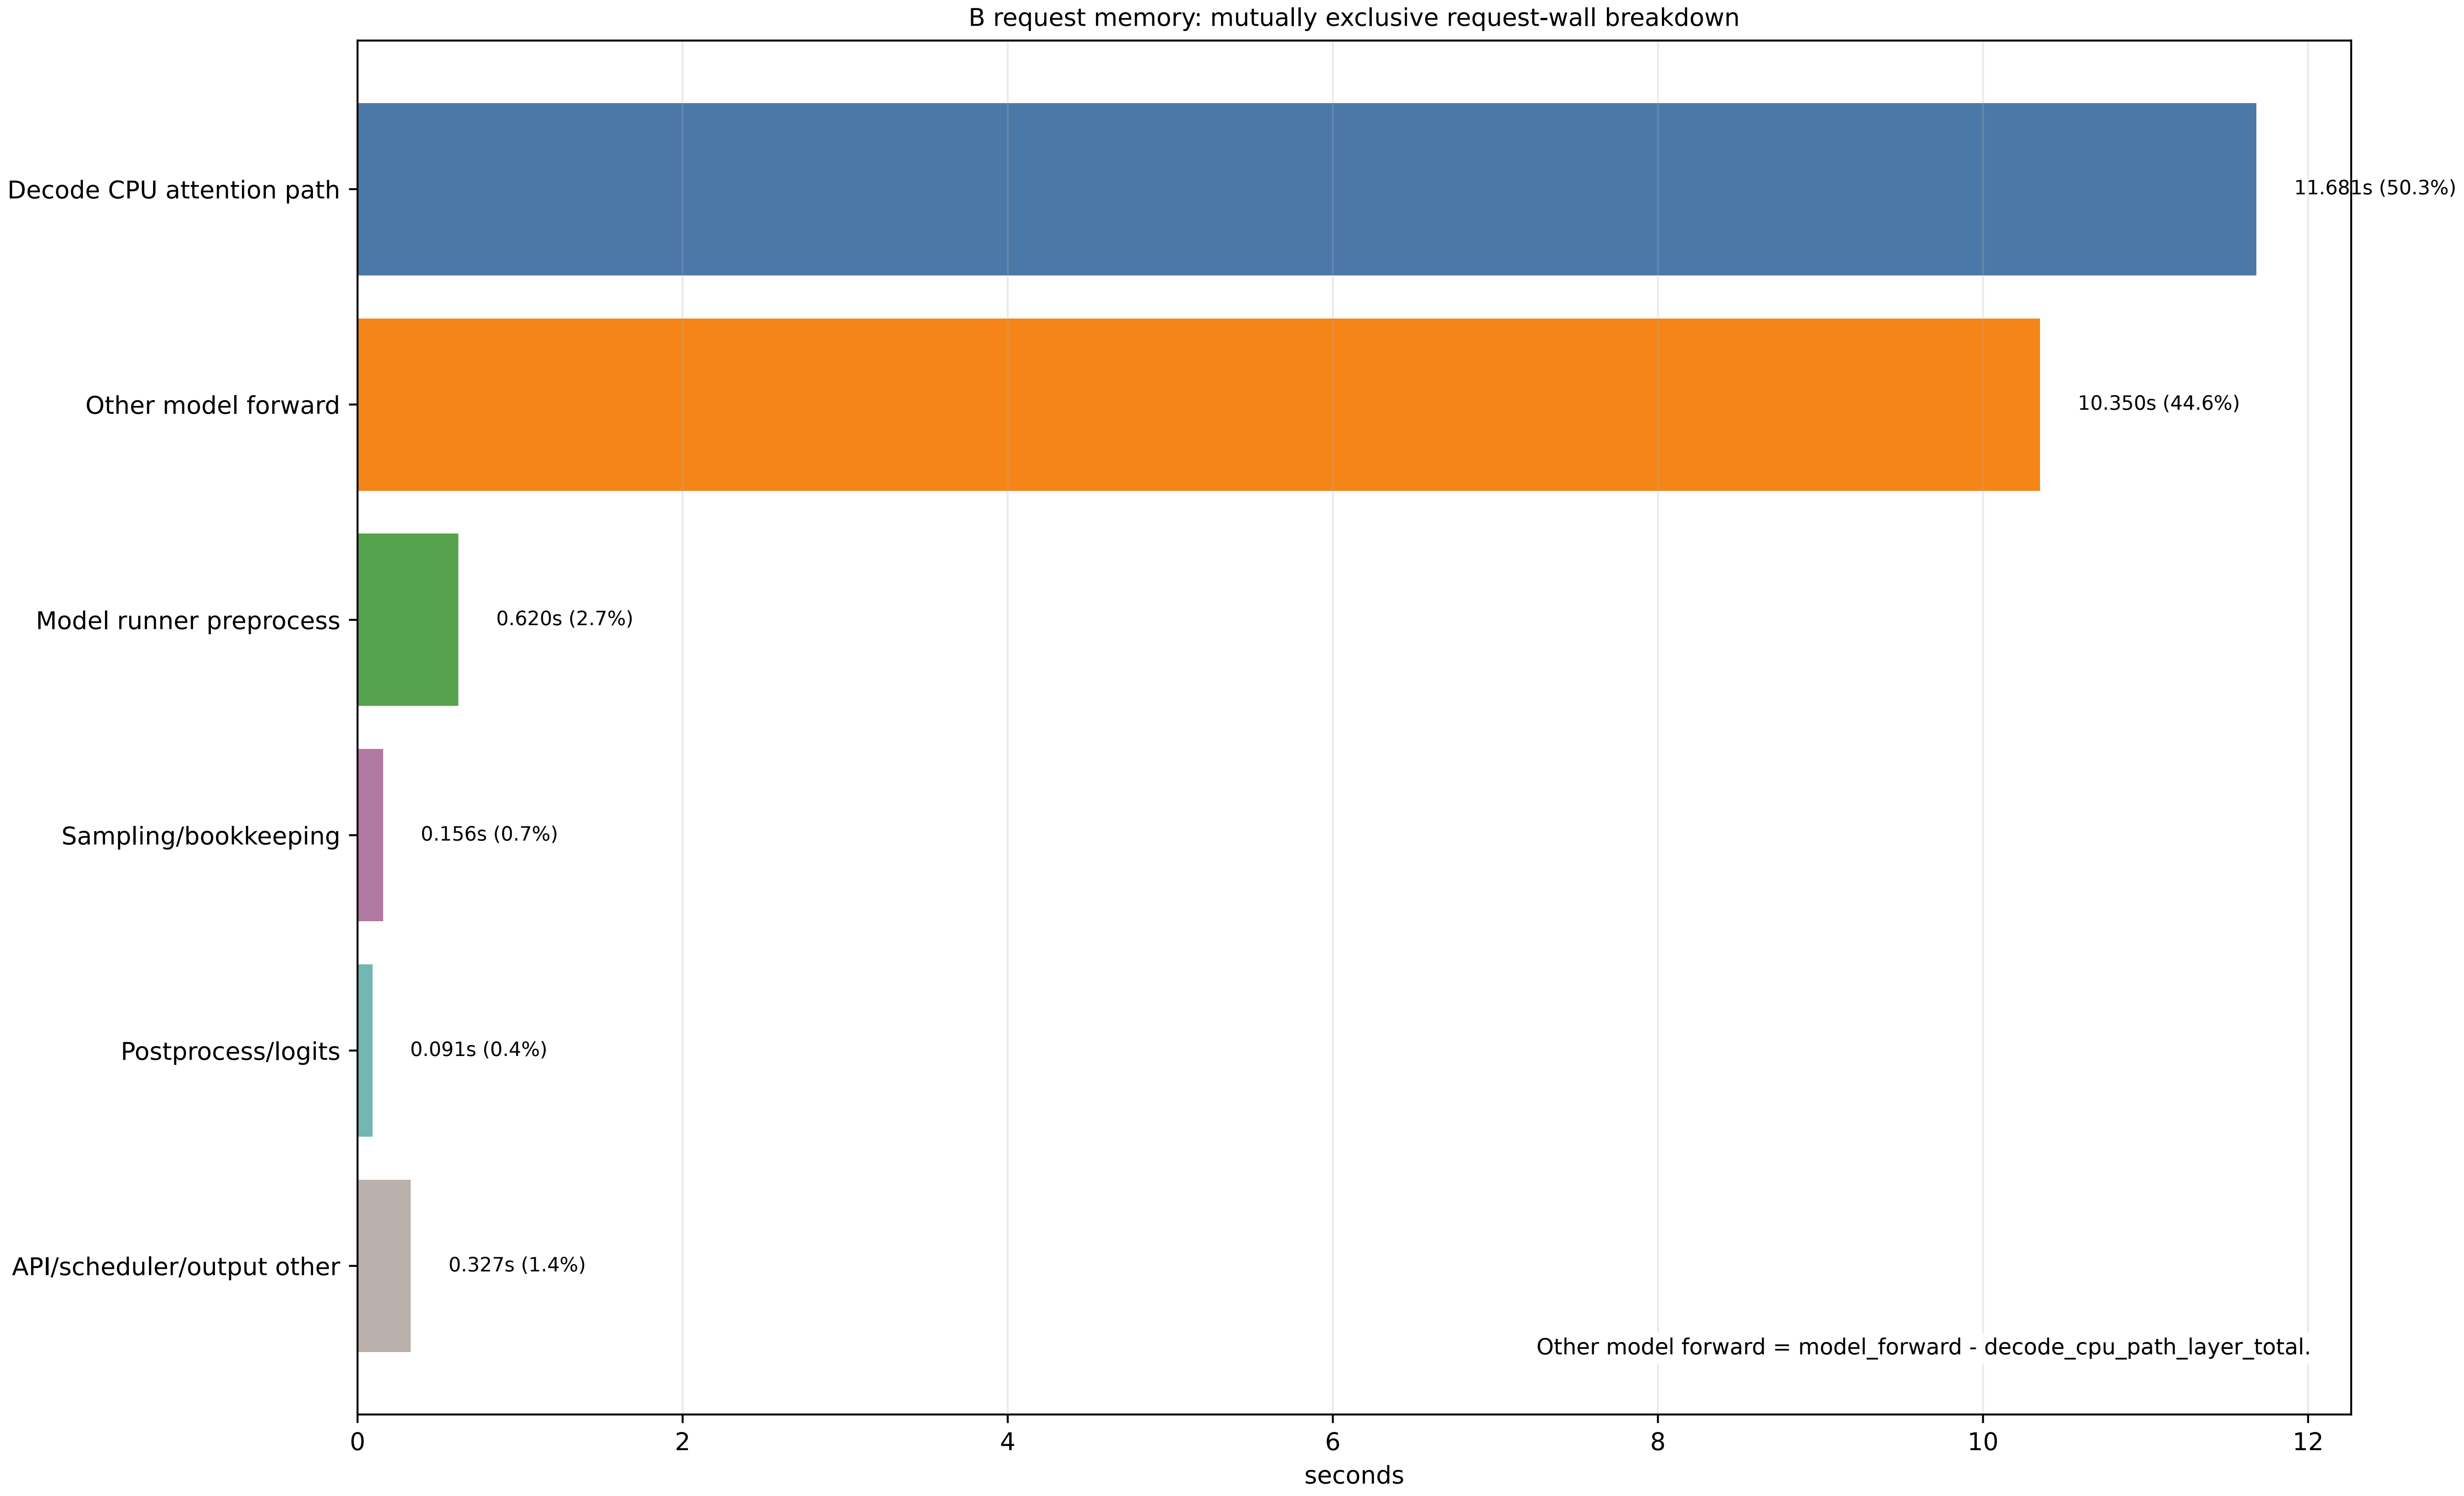
B request memory: mutually exclusive request-wall breakdown
Decode CPU attention path
Other model forward
Model runner preprocess
Sampling/bookkeeping
Postprocess/logits
API/scheduler/output other
11.681s (50.3%)
10.350s (44.6%)
0.620s (2.7%)
0.156s (0.7%)
0.091s (0.4%)
0.327s (1.4%)
Other model forward = model_forward - decode_cpu_path_layer_total.
0
2
4
6
8
10
12
seconds

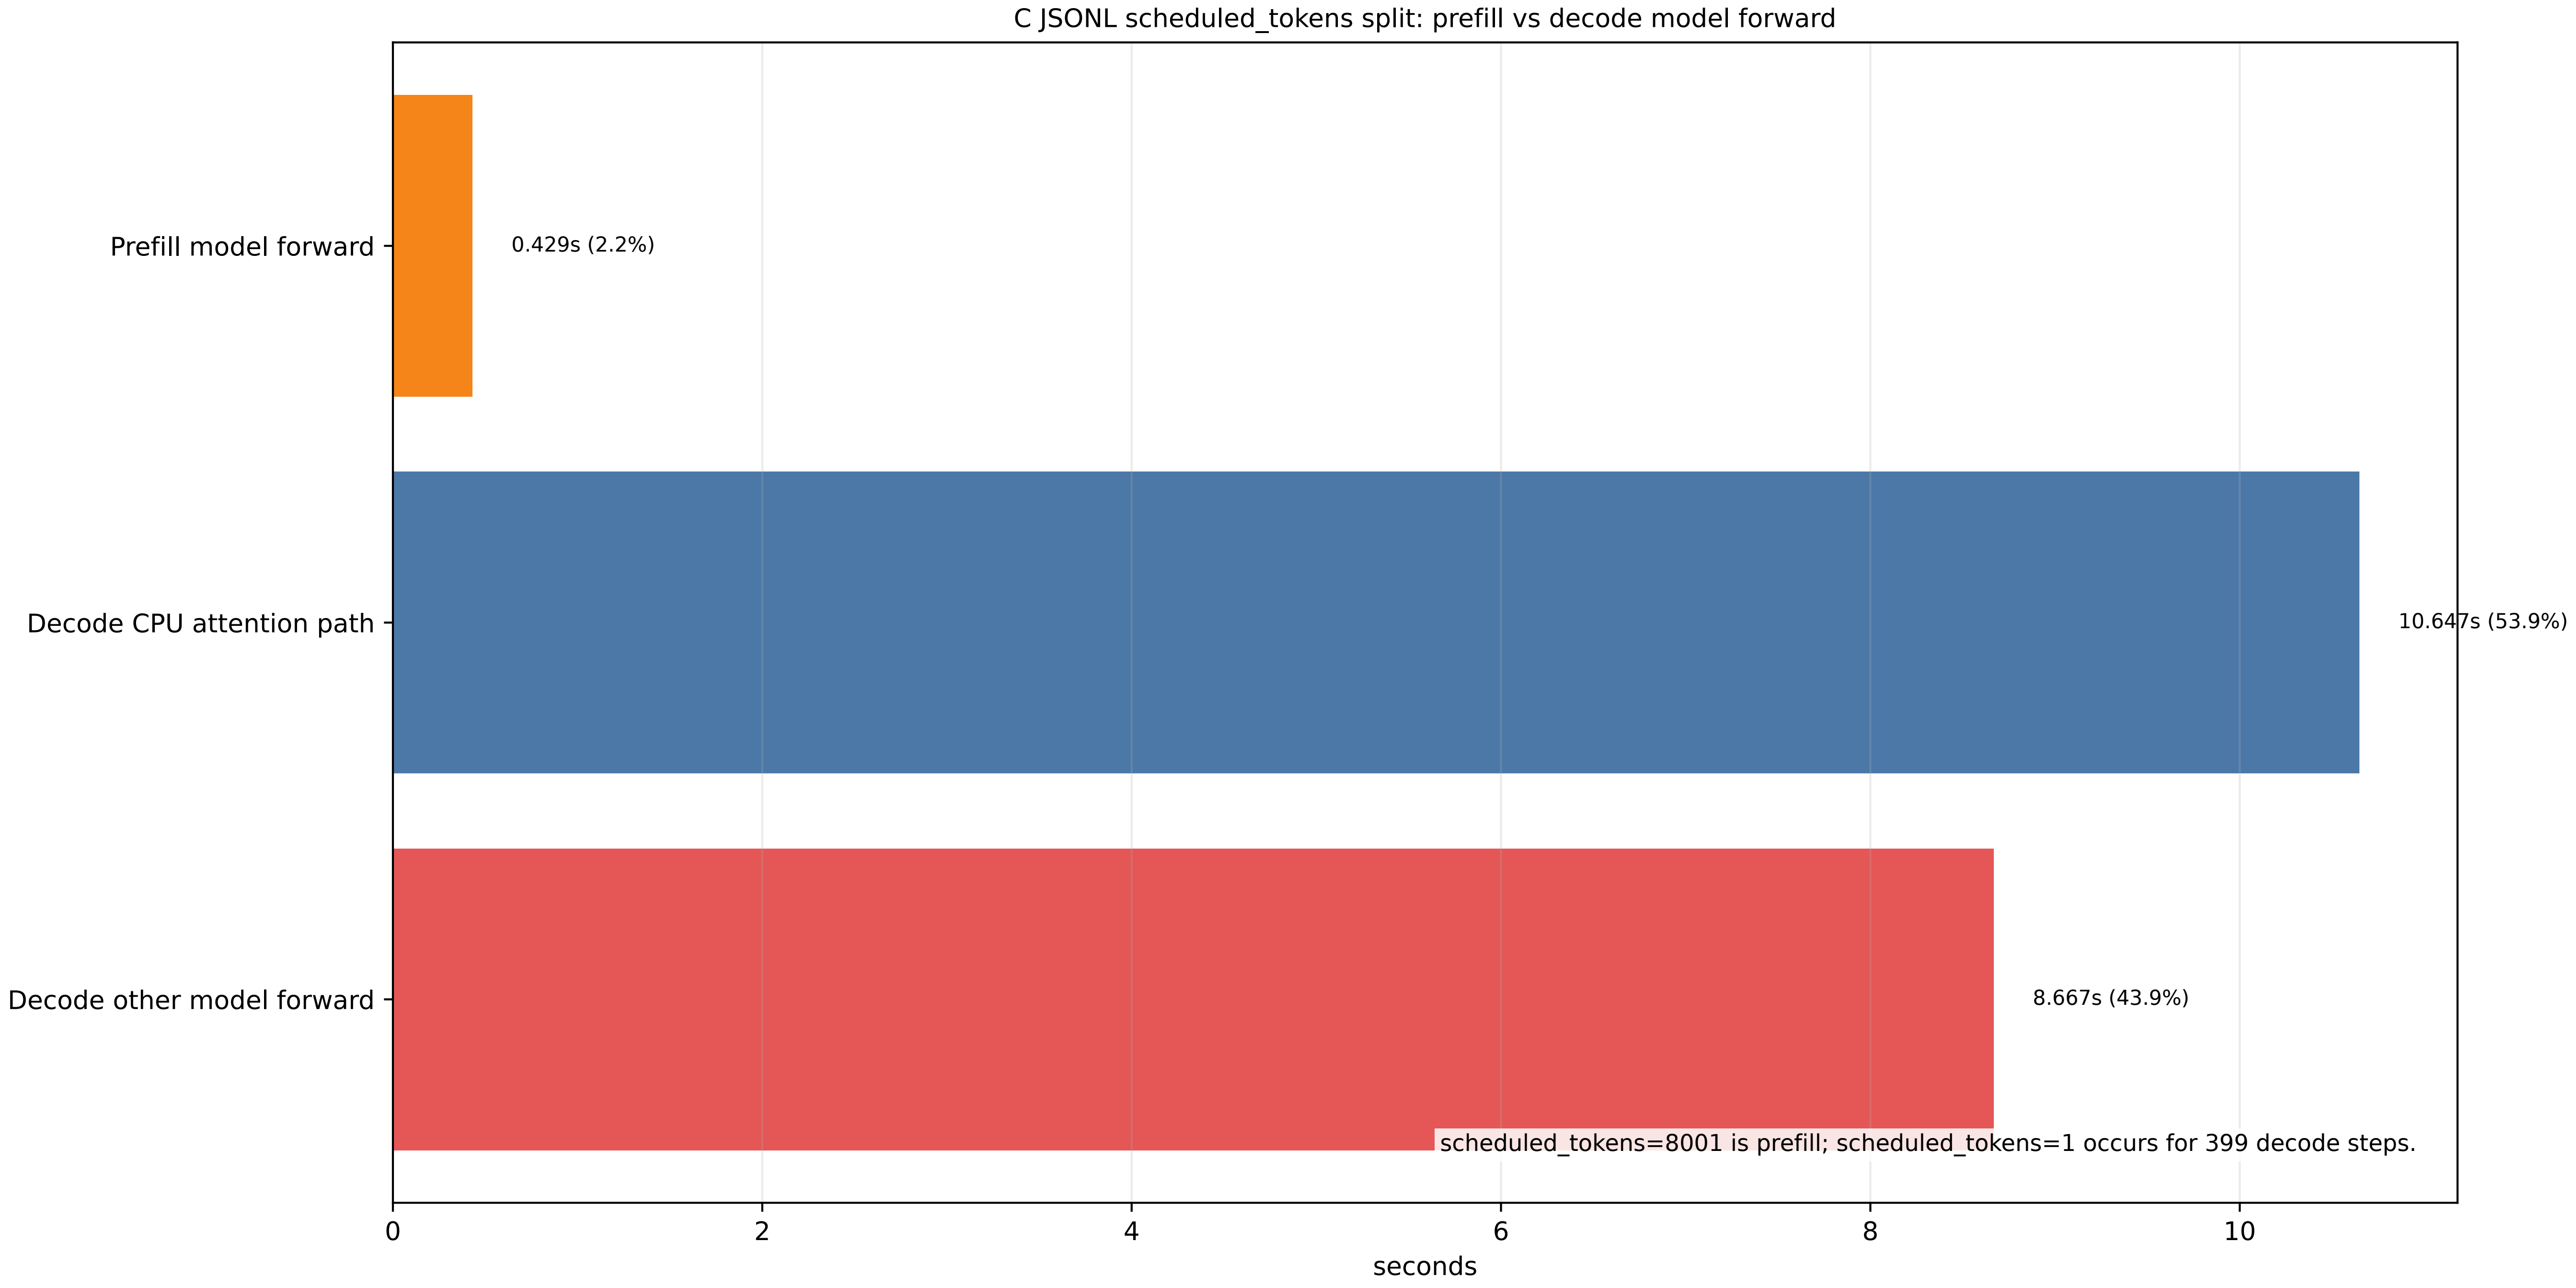
C JSONL scheduled_tokens split: prefill vs decode model forward
Prefill model forward
0.429s (2.2%)
Decode CPU attention path
10.647s (53.9%)
Decode other model forward
8.667s (43.9%)
scheduled_tokens=8001 is prefill; scheduled_tokens=1 occurs for 399 decode steps.
0
2
4
6
8
10
seconds

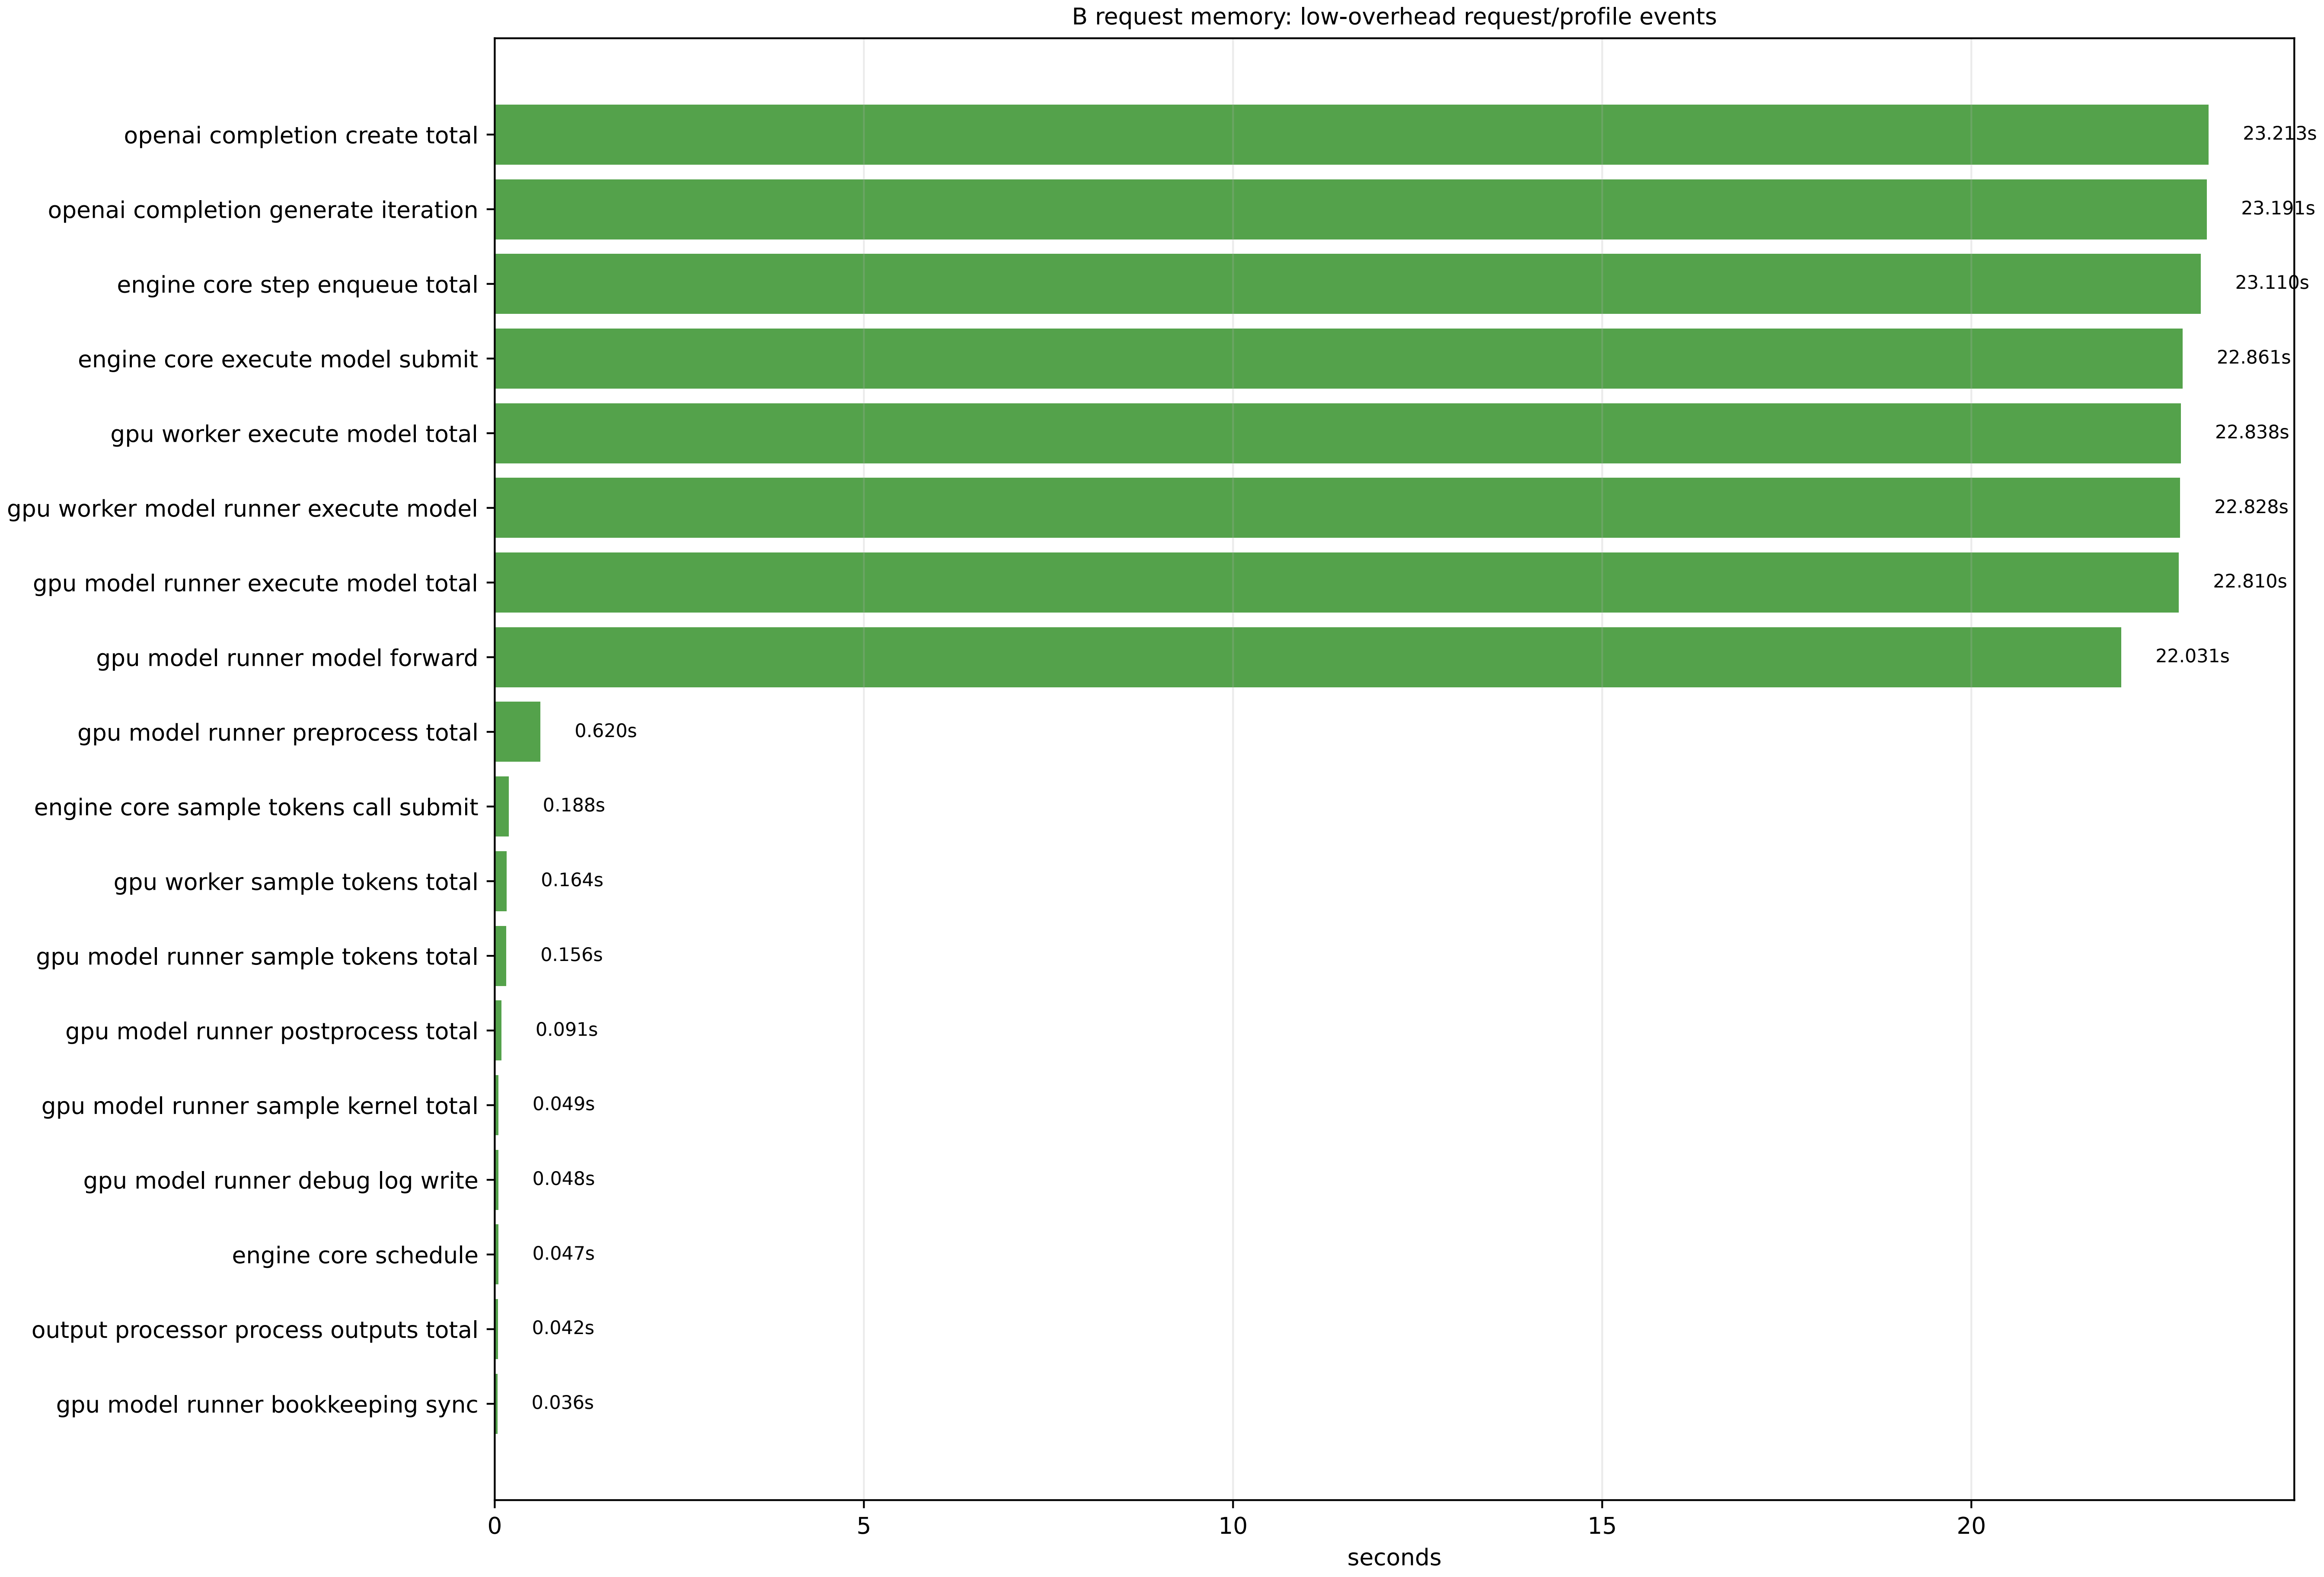
B request memory: low-overhead request/profile events
openai completion create total
openai completion generate iteration
engine core step enqueue total
engine core execute model submit
gpu worker execute model total
gpu worker model runner execute model
gpu model runner execute model total
gpu model runner model forward
gpu model runner preprocess total
engine core sample tokens call submit
gpu worker sample tokens total
gpu model runner sample tokens total
gpu model runner postprocess total
gpu model runner sample kernel total
gpu model runner debug log write
engine core schedule
output processor process outputs total
gpu model runner bookkeeping sync
23.213s
23.191s
23.110s
22.861s
22.838s
22.828s
22.810s
22.031s
0.620s
0.188s
0.164s
0.156s
0.091s
0.049s
0.048s
0.047s
0.042s
0.036s
0
5
10
15
20
seconds

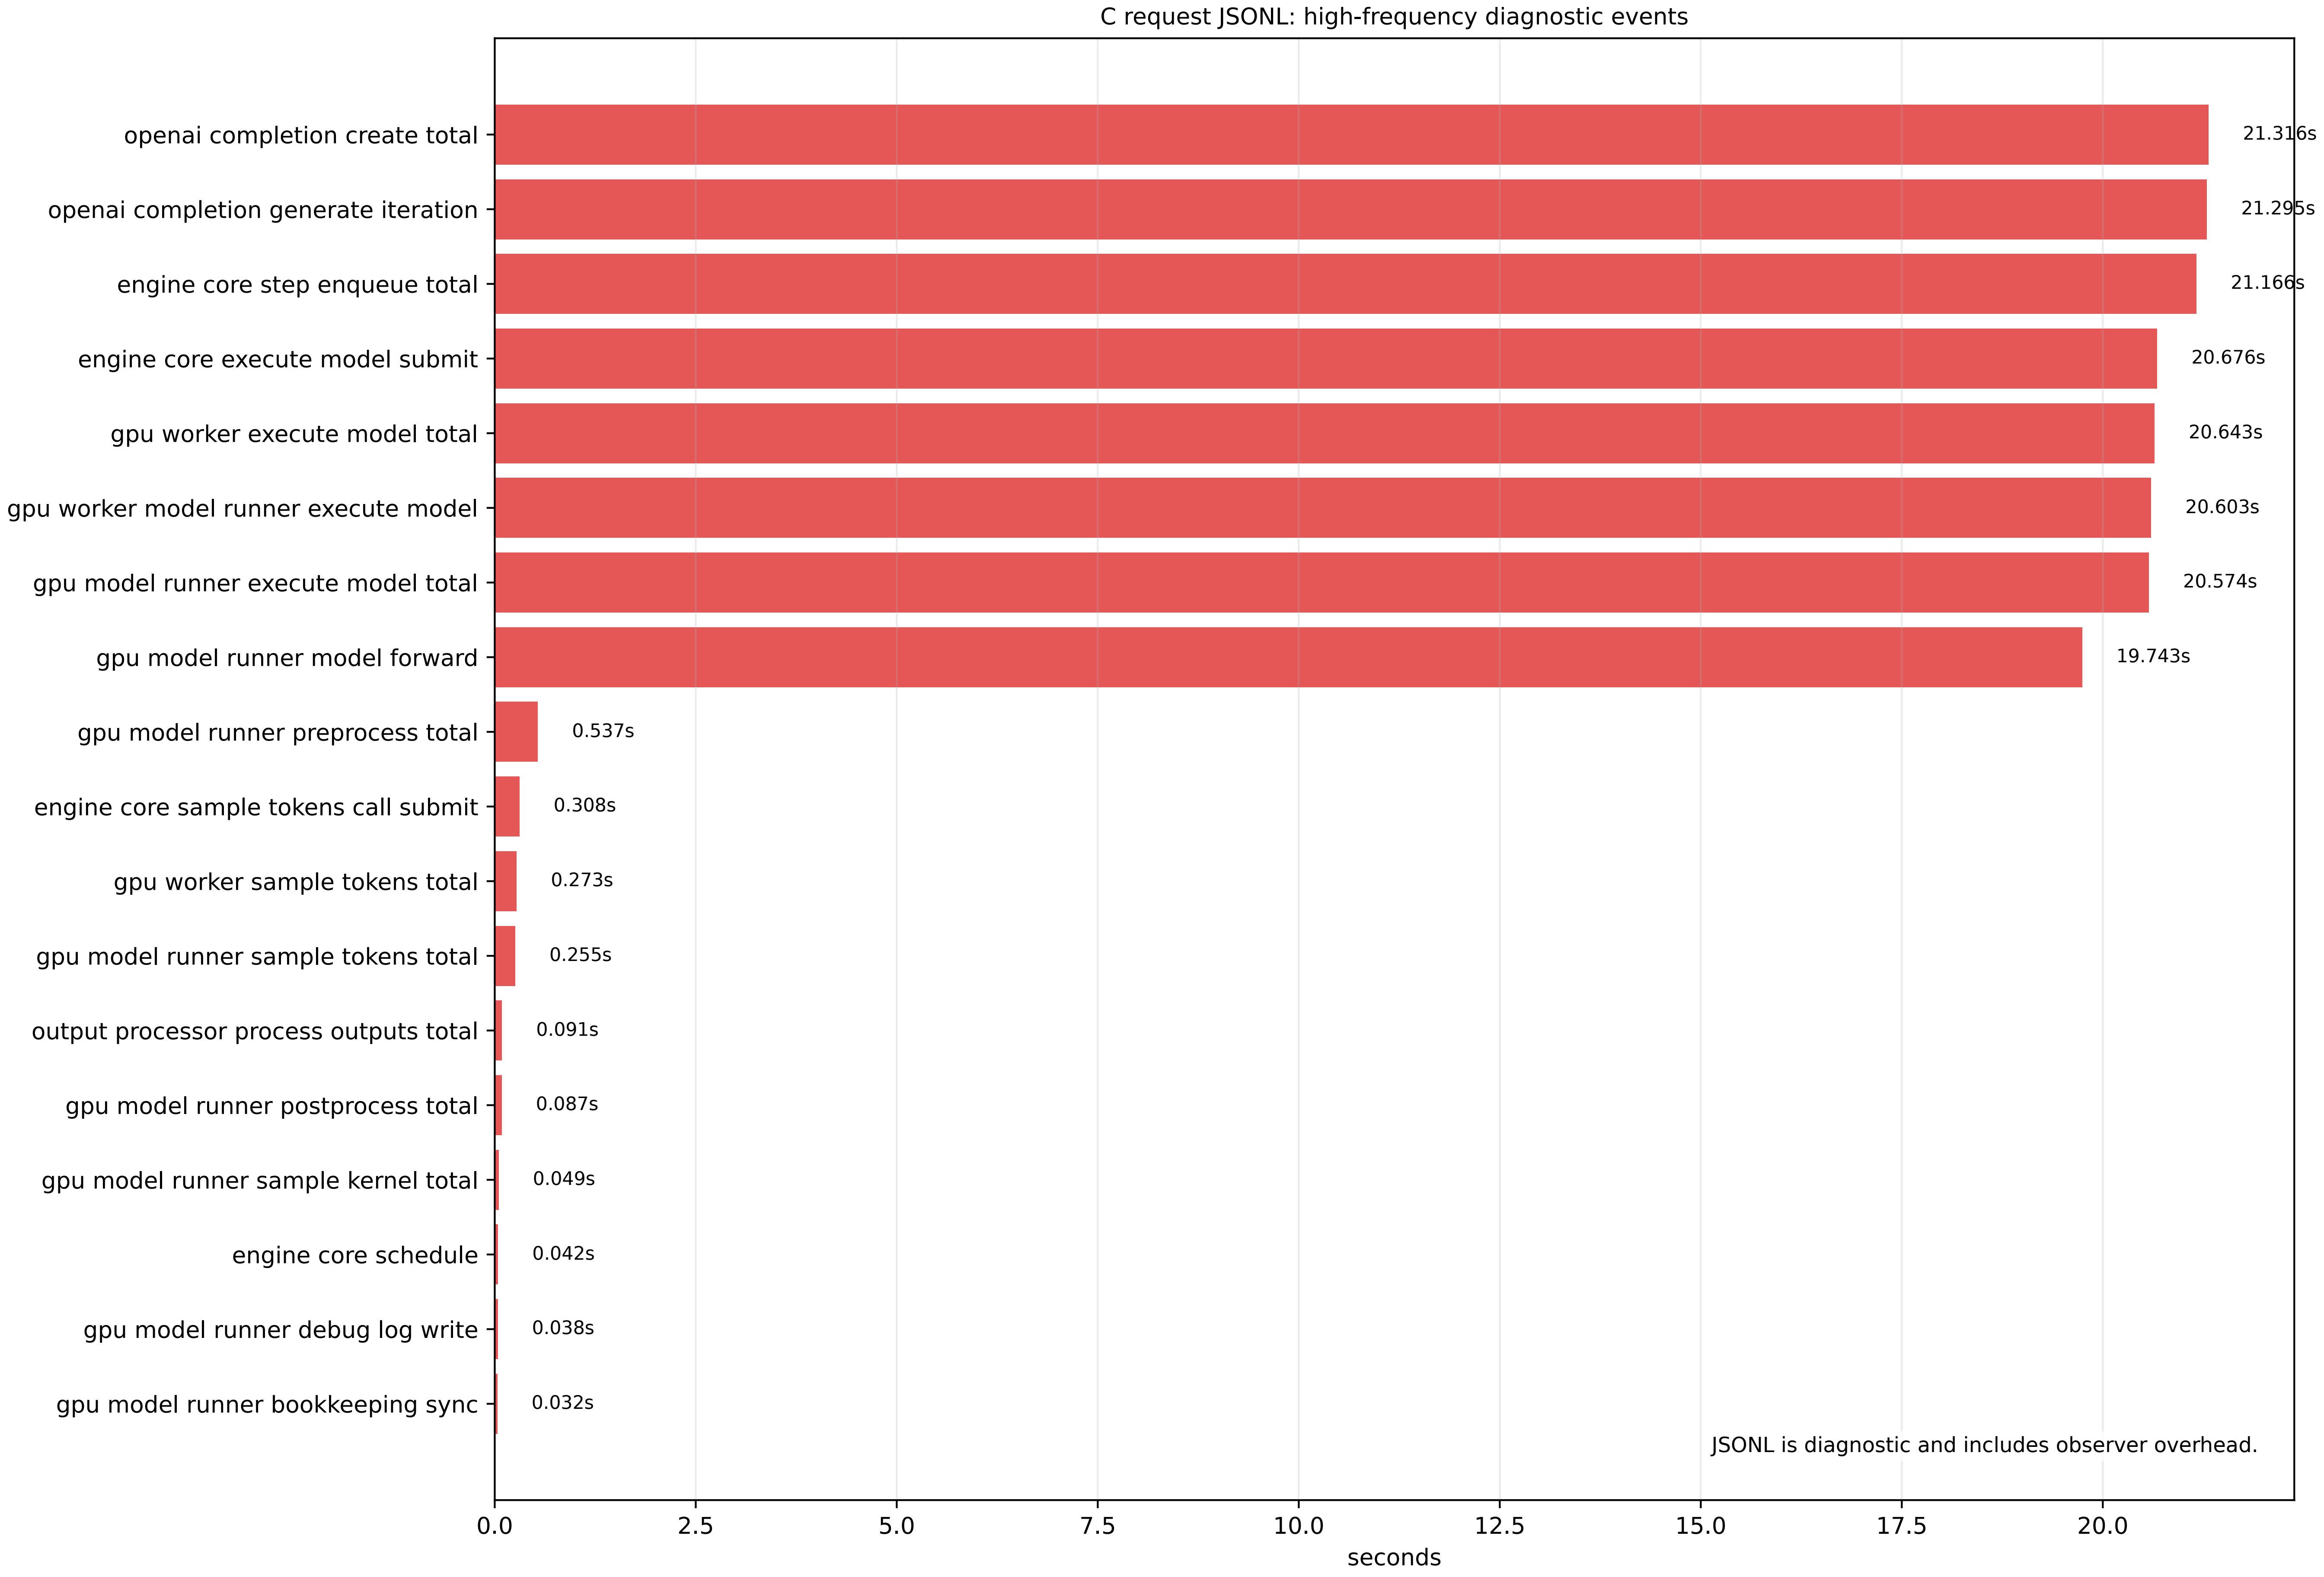
C request JSONL: high-frequency diagnostic events
openai completion create total
openai completion generate iteration
engine core step enqueue total
engine core execute model submit
gpu worker execute model total
gpu worker model runner execute model
gpu model runner execute model total
gpu model runner model forward
gpu model runner preprocess total
engine core sample tokens call submit
gpu worker sample tokens total
gpu model runner sample tokens total
output processor process outputs total
gpu model runner postprocess total
gpu model runner sample kernel total
engine core schedule
gpu model runner debug log write
gpu model runner bookkeeping sync
21.316s
21.295s
21.166s
20.676s
20.643s
20.603s
20.574s
19.743s
0.537s
0.308s
0.273s
0.255s
0.091s
0.087s
0.049s
0.042s
0.038s
0.032s
JSONL is diagnostic and includes observer overhead.
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
20.0
seconds

Generated: 2026-05-27 06:30:56 UTC

Workload: prompt=8000, output=400, batch=1, OMP_NUM_THREADS=120

12768 source: (completion_tokens - 1) * 32 = (400 - 1) * 32.

Baseline request wall - decode_cpu_path_layer_total = 10.403s.

Memory request profile observer delta vs baseline = 1.529s.

Memory profile request wall - decode_cpu_path_layer_total = 11.543s.

In B, other model forward after subtracting CPU attention = 10.350s.

Interpretation: use A/B for low-overhead timing; use C JSONL only to locate detailed events.